

Serialisation: a Comparison

Jeff Abrahamson

Jellybooks

Corba

Basics

- Backwards compatible
- Forwards compatible
- Efficiency (space, time)

Subtleties

- Community
- Safety
- Accent

Frameworks considered

Good:

- Thrift (Facebook, now Apache)
- Protobuf (Google)

Not so good:

- Boost Serialisation
- Json

Don't Even

- CORBA (1991) (cf. Vasa, 1628)
- XML-RPC (1998)
- SOAP (1998, successor to XML-RPC)

RPC

- Protobuf (no)
- Thrift (yes)
- Boost (no)
- Json (no)
- Cap'n Proto (yes)

Some popular RPC communication libraries:

- ZMQ
- RabbitMQ
- Etch (?)
- ZeroC ICE (?)
- Apache Qpid and AMQP (?)

But this isn't a talk about RPC.

Protobuf and Thrift

- IDL
- Proven scalable
- Community
- Similar efficiency

Protobuf

```
message Person {  
    required string name = 1;  
    required int32 id = 2;  
    optional string email = 3;  
  
    enum PhoneType {  
        MOBILE = 0;    // ...  
    }  
  
    message PhoneNumber {  
        required string number = 1;  
        optional PhoneType type = 2 [default = MOBILE];  
    }  
  
    repeated PhoneNumber phone = 4;  
}
```

Protobuf

```
message AddressBook {  
    repeated Person person = 1;  
}
```

Protobuf

```
// name
inline bool has_name() const;
inline void clear_name();
inline const ::std::string& name() const;
inline void set_name(const ::std::string& value);
inline void set_name(const char* value);
inline ::std::string* mutable_name();
```

Protobuf

//Checks if all the required fields have been set.

bool IsInitialized() **const**;

*// Returns a human-readable representation of the
// message, particularly useful for debugging.*

string DebugString() **const**;

*// Overwrites the message with the given message's
// values.*

void CopyFrom(**const** Person& from);

// Clears all the elements back to the empty state.

void Clear();

Protobuf

*// Serializes the message and stores the bytes in
// the given string. Note that the bytes are binary
// not text; we only use the string class as a
// convenient container.*

bool SerializeToString(string* output) **const**;

// Parses a message from the given string.

bool ParseFromString(**const** string& data);

// Writes the message to the given C++ ostream.

bool SerializeToOstream(ostream* output) **const**;

// Parses a message from the given C++ istream.

bool ParseFromIstream(istream* input);

Protobuf

```
AddressBook address_book;
```

```
Person* person = address_book.add_person();  
person->set_id(id);  
person->set_name(name);  
if (!email.empty()) {  
    person->set_email(email);  
}
```

```
Person::PhoneNumber* phone_number  
    = person->add_phone();  
phone_number->set_number(number);  
phone_number->set_type(Person::MOBILE);
```

Protobuf

```
for (int j = 0; j < person.phone_size(); j++) {  
    const Person::PhoneNumber&  
        phone_number = person.phone(j);
```



```
const map<string , string > MAPCONSTANT =  
    { 'hello ' : 'world ',  
      'goodnight ' : 'moon' }
```

```
enum Operation {  
    ADD = 1,  
    // ...  
}
```

```
struct Work {  
    1: i32 num1 = 0,  
    2: i32 num2,  
    3: Operation op,  
    4: optional string comment,  
}
```

```
exception InvalidOperation {  
    1: i32 whatOp,  
    2: string why  
}
```

```
struct SharedStruct {  
    1: i32 key  
    2: string value  
}
```

```
service SharedService {  
    SharedStruct getStruct(1: i32 key)  
}
```

```
service Calculator extends shared.SharedService {  
    void ping(),  
    i32 add(1:i32 num1, 2:i32 num2),  
    i32 calculate(1:i32 logid, 2:Work w)  
        throws (1:InvalidOperation ouch),  
  
    /**  
     * This method has a oneway modifier. That  
     * means the client only makes a request and  
     * does not listen for any response at  
     * all. Oneway methods must be void.  
     */  
    oneway void zip()  
}
```

Thrift

```
boost::shared_ptr<TTransport> socket(  
    new TSocket("localhost", 9090));  
boost::shared_ptr<TTransport> transport(  
    new TBufferedTransport(socket));  
boost::shared_ptr<TProtocol> protocol(  
    new TBinaryProtocol(transport));  
CalculatorClient client(protocol);
```

Thrift

```
transport->open();
```

```
client.ping();
```

```
cout << "1 + 1 = " << client.add(1, 1) << endl;
```

```
Work work;  
work.op = Operation::DIVIDE;  
work.num1 = 1;  
work.num2 = 0;  
  
try {  
    client.calculate(1, work);  
    cout << "Division by zero!" << endl;  
} catch (InvalidOperation& io) {  
    cout << "InvalidOperation: " << io.why << endl;  
    cout << io << endl;    // Same thing.  
}
```

Thrift

```
SharedStruct ss;  
client.getStruct(ss, 1);  
cout << "Received log: " << ss << endl;  
  
transport->close();
```

```
int main() {  
    boost::shared_ptr<TProtocolFactory>  
        protocolFactory(  
            new TBinaryProtocolFactory());  
    boost::shared_ptr<CalculatorHandler>  
        handler(new CalculatorHandler());  
    boost::shared_ptr<TProcessor> processor(  
        new CalculatorProcessor(handler));  
    boost::shared_ptr<TServerTransport>  
        serverTransport(new TServerSocket(9090));  
    boost::shared_ptr<TTransportFactory>  
        transportFactory(  
            new TBufferedTransportFactory());  
  
    TSimpleServer server(processor, serverTransport,  
                        transportFactory,  
                        protocolFactory);  
}
```



```
// Server
```

```
int32_t add(const int32_t n1,  
           const int32_t n2) { ...
```

```
int32_t calculate(const int32_t logid,  
                 const Work& work) { ...
```

```
void getStruct(SharedStruct& ret,  
              const int32_t logid) { ...
```

```
void zip() { ...
```

Thrift

```
boost::shared_ptr<TFileTransport>  
    transport(new TFileTransport(filename));  
boost::shared_ptr<TBinaryProtocol>  
    protocol(new TBinaryProtocol(transport));  
myObj.write(protocol.get());
```

Boost

// Save: << or the & operator

ar << data;

ar & data;

// Load: >> or the & operator.

ar >> data;

ar & data;

```
// include headers that implement an archive in  
// simple text format  
#include <boost/archive/text_oarchive.hpp>  
#include <boost/archive/text_iarchive.hpp>
```

```
class gps_position {  
    private:  
        friend class boost::serialization::access;  
        // When the class Archive corresponds to an  
        // output archive, the & operator is defined  
        // similar to <<. Likewise, when the class  
        // Archive is a type of input archive the &  
        // operator is defined similar to >>.  
        template<class Archive>  
        void serialize(Archive & ar,  
                      const unsigned int version) {  
            ar & degrees;  
            ar & minutes;  
            ar & seconds;  
        }  
        int degrees, minutes;  
        float seconds;
```

Boost

```
int main() {  
    std::ofstream ofs("filename");  
    const gps_position g(35, 59, 24.567f);  
  
    {  
        boost::archive::text_oarchive oa(ofs);  
        oa << g;  
    }  
  
    gps_position newg;    // Later...  
    {  
        std::ifstream ifs("filename");  
        boost::archive::text_iarchive ia(ifs);  
        ia >> newg;  
    }  
    return 0;  
}
```

```
// Non-intrusive version.  
template<class Archive>  
void serialize(Archive & ar, gps_position & g,  
               const unsigned int version) {  
    ar & g.degrees;  
    ar & g.minutes;  
    ar & g.seconds;  
}
```

```
// Serializable members.  
class bus_stop {  
    friend class boost::serialization::access;  
    template<class Archive>  
    void serialize(Archive & ar,  
                  const unsigned int version) {  
        ar & latitude;  
        ar & longitude;  
    }  
    gps_position latitude;  
    gps_position longitude;  
};
```



```
// Derived classes.
#include <boost/serialization/base_object.hpp>

class bus_stop_corner : public bus_stop {
    friend class boost::serialization::access;
    template<class Archive>
    void serialize(Archive & ar,
                   const unsigned int version) {
        // serialize base class information
        ar & boost::serialization::base_object<
            bus_stop>(*this);
        ar & street1;
        ar & street2;
    }
    std::string street1;
    std::string street2;
};
```

```
// Pointers.  
class bus_route {  
    friend class boost::serialization::access;  
    bus_stop* stops[10];  
    template<class Archive>  
    void serialize(Archive & ar,  
                   const unsigned int version)  
{  
    for(int i = 0; i < 10; ++i) {  
        ar & stops[i];  
    }  
}
```

```
// Arrays.  
class bus_route {  
    friend class boost::serialization::access;  
    bus_stop* stops[10];  
    template<class Archive>  
    void serialize(Archive & ar,  
                   const unsigned int version) {  
        ar & stops;  
    }  
}
```

```
// STL.  
#include <boost/serialization/list.hpp>  
  
class bus_route {  
    friend class boost::serialization::access;  
    std::list<bus_stop*> stops;  
    template<class Archive>  
    void serialize(Archive & ar,  
                  const unsigned int version) {  
        ar & stops;  
    }  
}
```

```
// Versioning (1).
#include <boost/serialization/list.hpp>
#include <boost/serialization/string.hpp>

class bus_route {
    friend class boost::serialization::access;
    std::list<bus_stop*> stops;
    std::string driver_name;
    template<class Archive>
    void serialize(Archive & ar,
                  const unsigned int version) {
        ar & driver_name;    // !!
        ar & stops;
    }
}
```

```
// Versioning (2).
#include <boost/serialization/list.hpp>
#include <boost/serialization/string.hpp>
#include <boost/serialization/version.hpp>

class bus_route {
    friend class boost::serialization::access;
    std::list<bus_stop *> stops;
    std::string driver_name;
    template<class Archive>
    void serialize(Archive & ar,
                   const unsigned int version) {
        if (version > 0)
            ar & driver_name;
        ar & stops;
    }
}
```

```
BOOST_CLASS_VERSION( bus_route , 1)
```

Json

- Ubiquitous
- Human readable
- No validation or error checking
- Verbose

Cf. Bson, MessagePack, ...

Questions?

`https://github.com/JeffAbrahamson/talks/`

`jeff@purple.com`